



T-Fuzz: Fuzzing by Program Transformation

Hui Peng¹, Yan Shoshitaishvili², Mathias Payer¹



Fuzzing as a bug finding approach

- Fuzzing is finding more and more CVEs
- Vendors use it as proactive defense measure: OSS-Fuzz
- Hackers use it as first step in exploit development



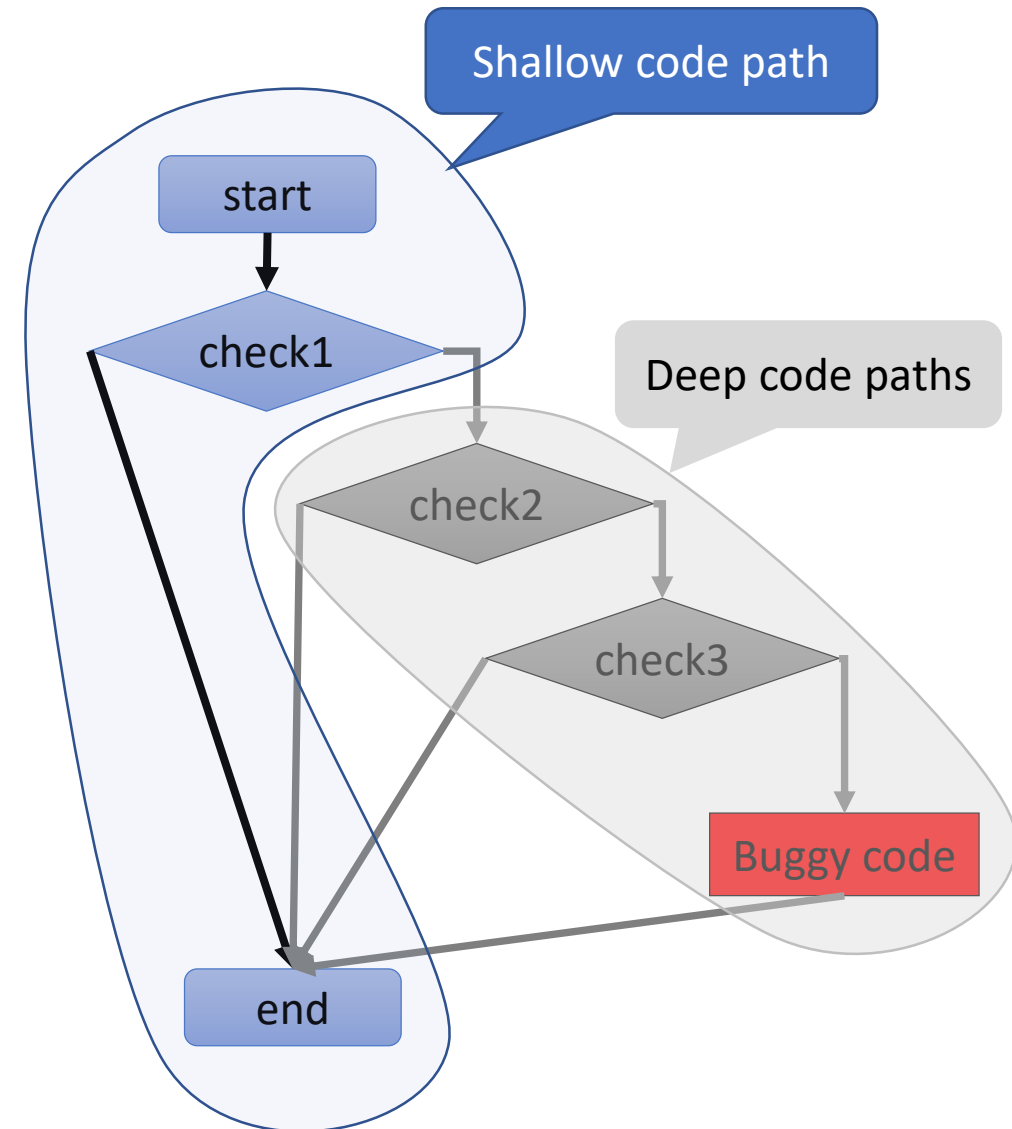
Challenges for fuzzers

➤ Challenges

- ❑ Shallow coverage
- ❑ Hard to find “deep” bugs

➤ Root Cause

- ❑ Fuzzer-generated inputs cannot bypass complex sanity checks in the target program



Existing approaches & their limitations

➤ Existing approaches focus on *input generation*

- ❑ Driller (concolic execution)
- ❑ VUzzer (taint analysis, data & control flow analysis)

➤ Limitations

- ❑ High overhead
- ❑ Not scalable
- ❑ Not able to bypass “hard” checks
 - Checks on checksum, crypto-hash values

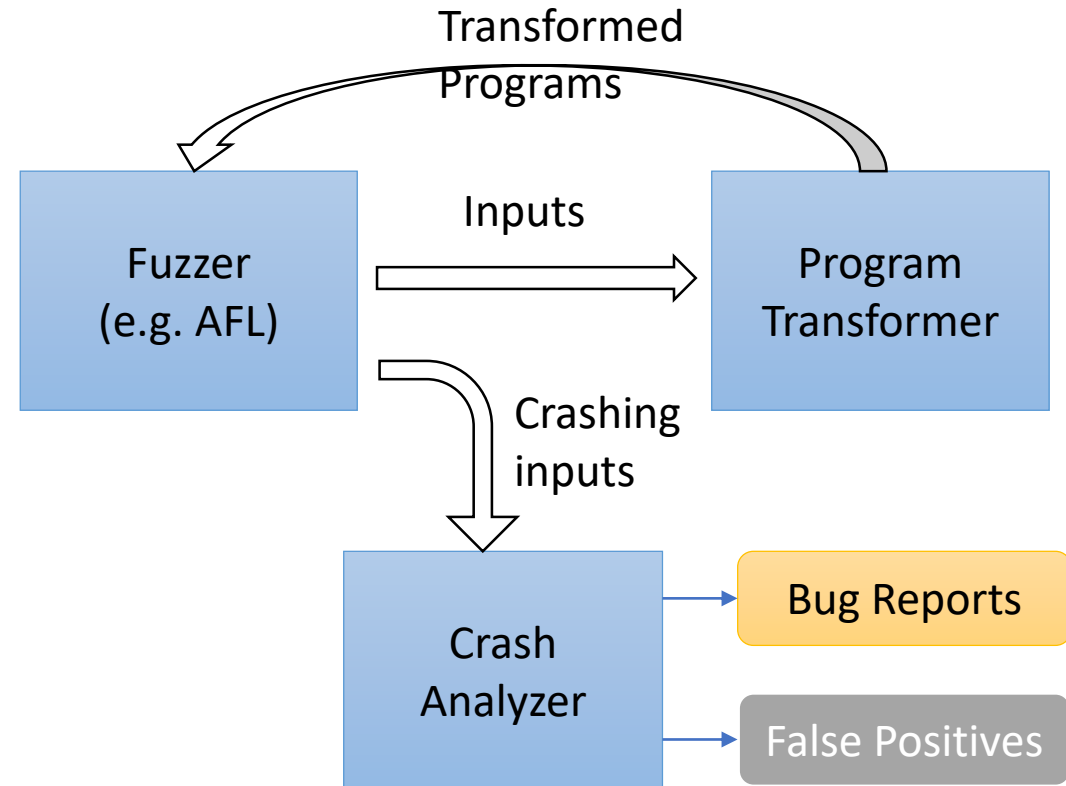
Insight: some checks are non-critical

- Some sanity checks are not intended to prevent bugs
- **Non-Critical Checks (NCC)**
 - ❑ E.g., check on magic values, checksums, hashes
- Removing NCCs won't incur erroneous bugs
- Removal of NCCs simplifies fuzzing

```
void main() {  
    int fd = open(...);  
    char *hdr = read_header(fd);  
    if (strncmp(hdr, "ELF", 3) == 0) {  
        // main program logic  
        // ...  
    } else {  
        error();  
    }  
}
```

T-Fuzz: fuzzing by program transformation


- Fuzzer generates inputs
- When Fuzzer gets stuck,
Program Transformer:
 - ❑ Detects **NCC candidates**
 - ❑ Transforms program
- Crash Analyzer verifies crashes
- Repeat



T-Fuzz design

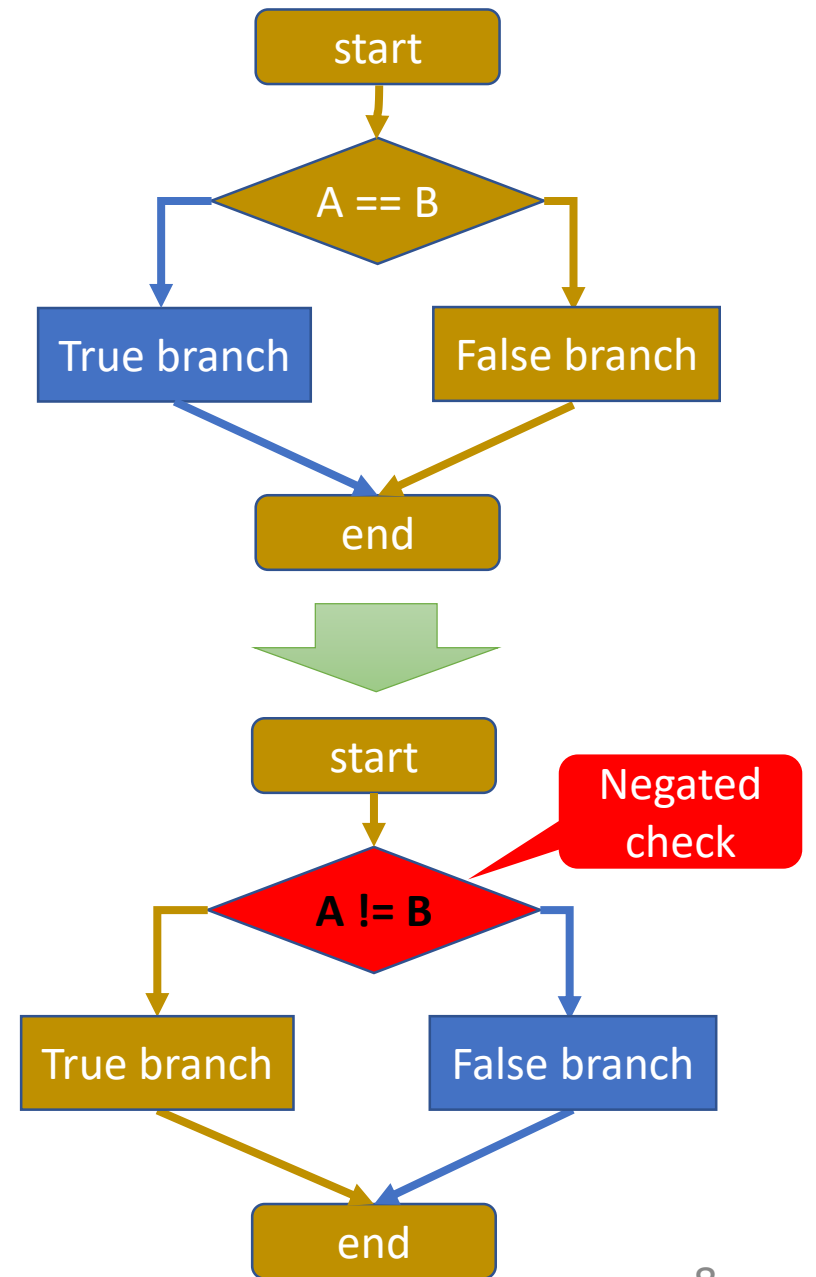
Detecting NCC candidates

- Approximate NCCs as the edges connecting covered/uncovered nodes in the CFG
- Overapproximate, may contain false positive
- Lightweight and simple to implement
 - ❑ dynamic tracing

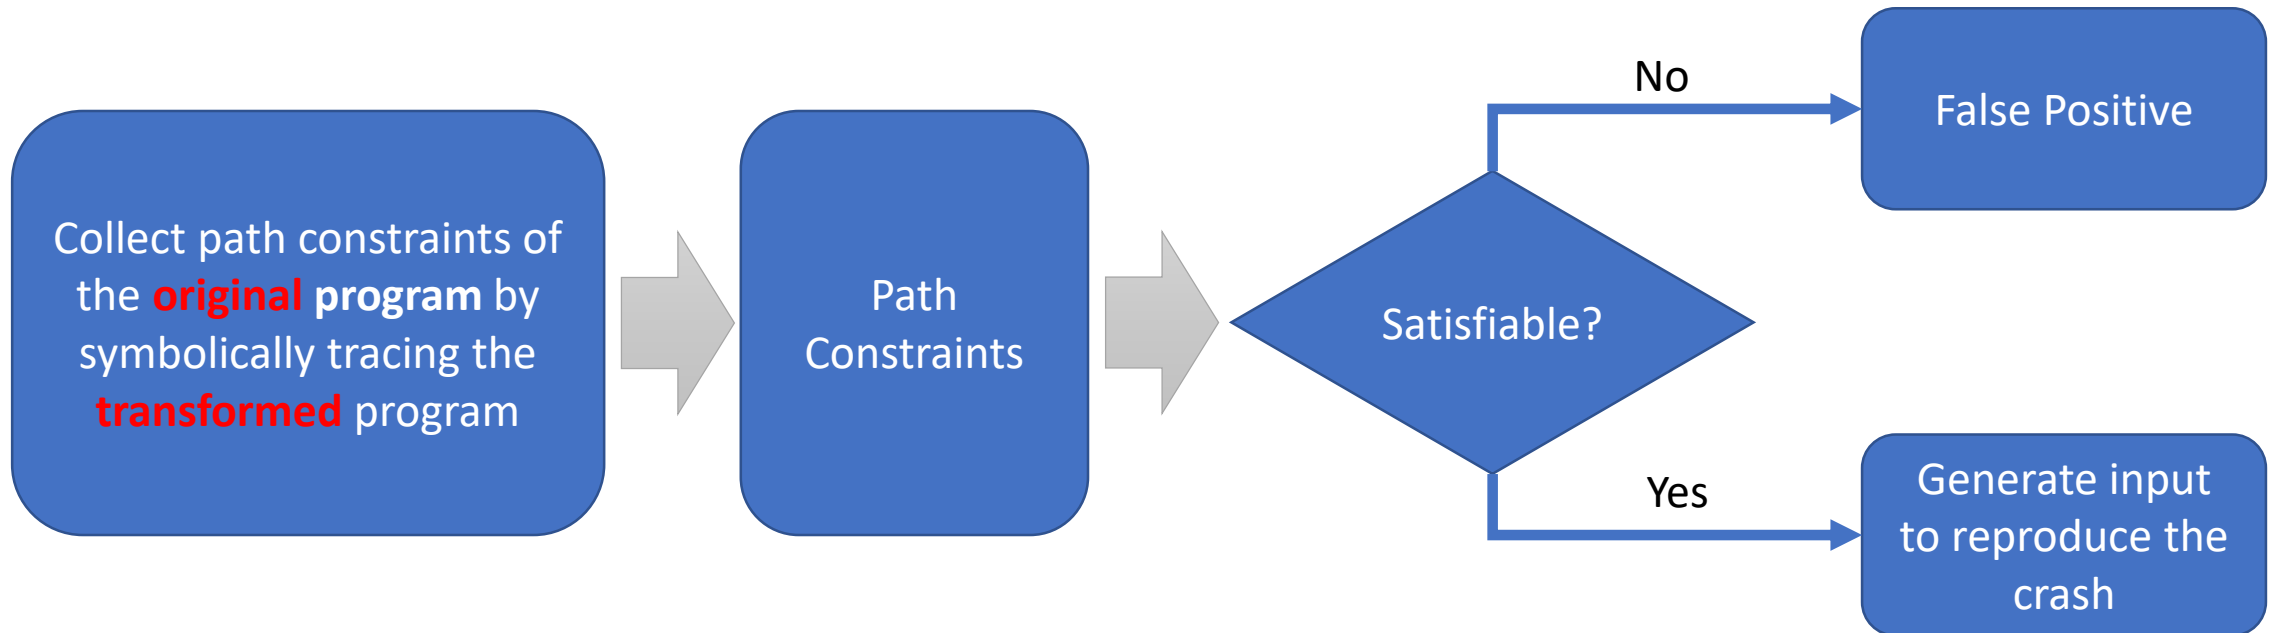


Program Transformation

- **Goal:** disable NCCs
- Our approach: **negate NCCs**
 - ❑ Easy to implement: static binary rewriting
 - ❑ Zero runtime overhead in target program
 - ❑ The CFG of the program stays the same
 - ❑ Traces of the transformed program map to the original one
 - ❑ Path constraints of the original program can be recovered



Filtering out false positives & reproducing bugs



Example 1

```
int main () {  
    int x = read_input();  
    int y = read_input();  
    if (x > 0) {  
        if (y == 0xdeadbeef)  
            bug();  
    }  
}
```

Original Program



```
int main () {  
    int x = read_input();  
    int y = read_input();  
    if (x > 0) {  
        if (y != 0xdeadbeef)  
            bug();  
    }  
}
```

Transformed Program

Collected path constraints:

{x > 0, y == 0xdeadbeef}

SAT

True BUG

Un-negating

Negated check

Example 2

UNSAT

False BUG

```
int main () {  
    int i = read_input();  
  
    if (i > 0) {  
        func(i);  
    }  
}  
  
void func(int i) {  
    if (i <= 0) {  
        bug();  
    }  
    //...  
}
```

Original Program



```
int main () {  
    int i = read_input();  
  
    if (i > 0) {  
        func(i);  
    }  
}  
  
void func(int i) {  
    if (i > 0) {  
        bug();  
    }  
    //...  
}
```

Transformed Program

path constraints:

{i > 0; i <= 0}

Un-negating

Negated check

Limitations of T-Fuzz (1)

- False crashes may hinder discovery of true bugs (L1)

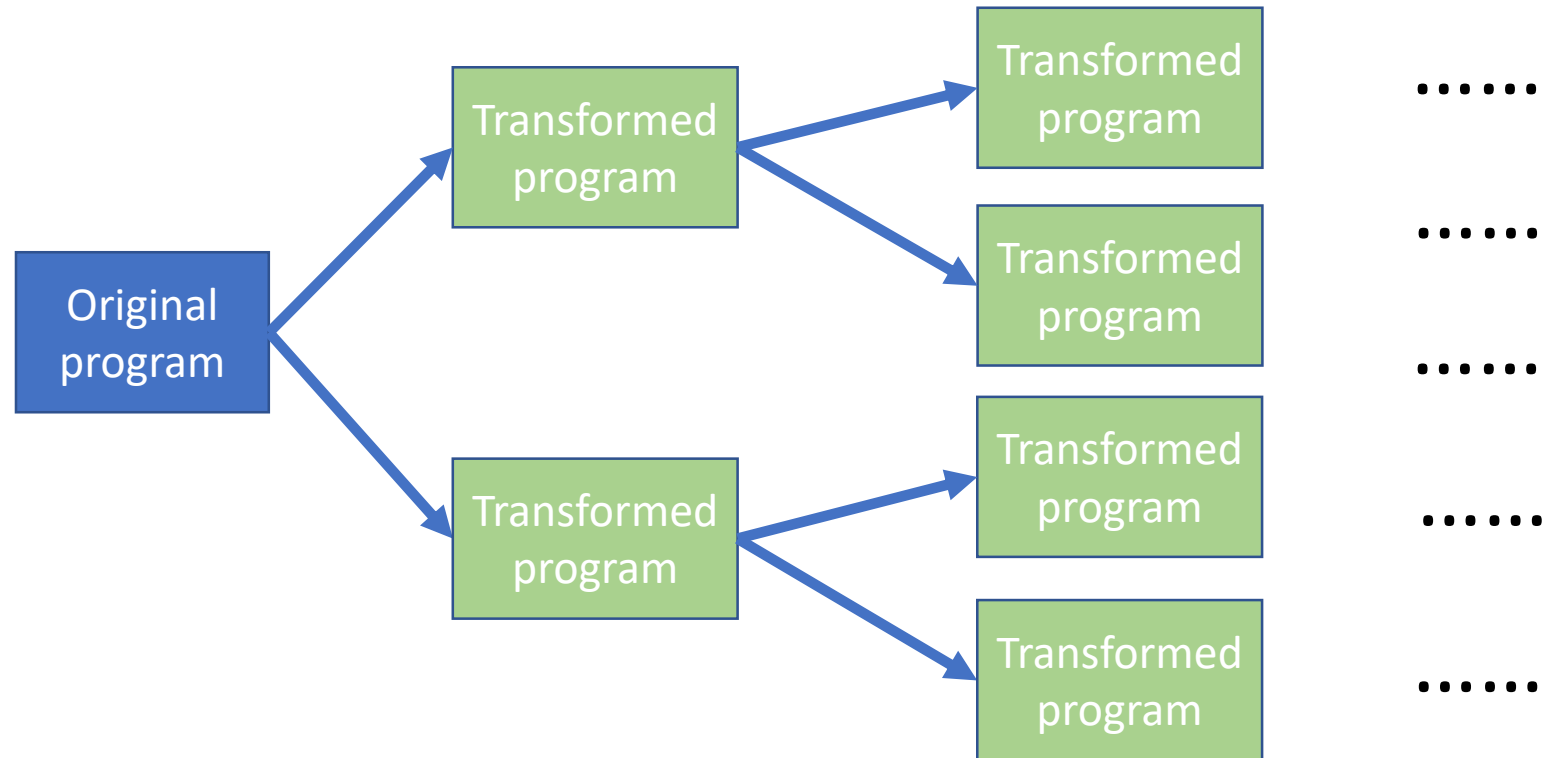
```
FILE *fp = fopen(...);  
if (fp != NULL) {  
    // False crash  
    fread(fp, ...);  
    // ...  
    // true bug  
    bug();  
}
```

Example: false crash hindering discovery of true bug

Limitations of T-Fuzz (2)

➤ Transformation explosion (L2)

➤ Analogous to path explosion issue in symbolic execution

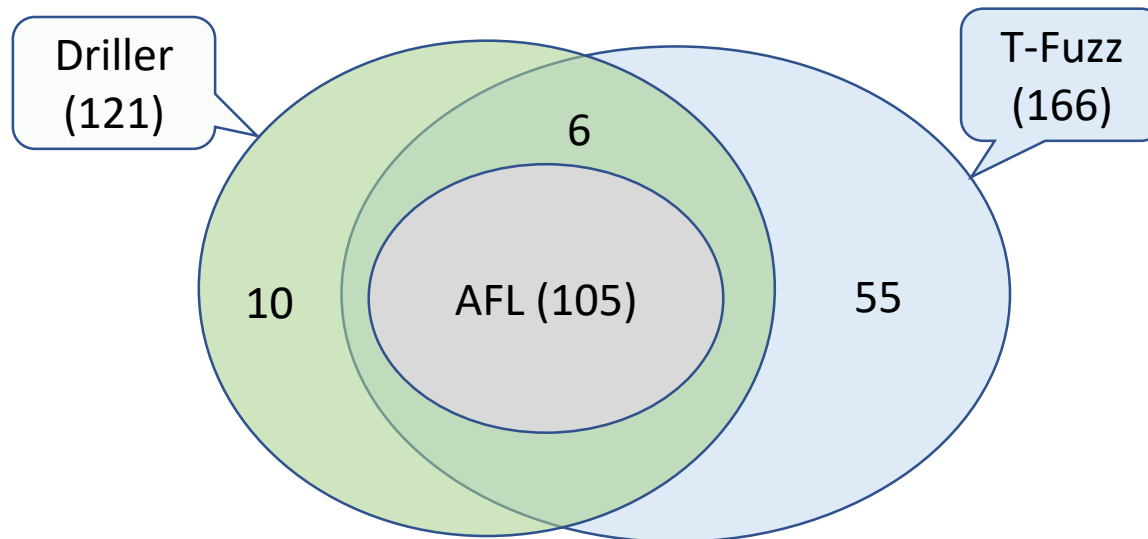


Evaluation

- DARPA CGC dataset
- LAVA-M dataset
- 4 real-world programs

CGC dataset

- Improvement over Driller/AFL: **55 (45%)/61 (58%)**
- T-Fuzz is defeated by Driller in 10
 - ❑ due to false crashes (L1) in 3
 - ❑ due to transformation explosion (L2) in 7



Method	# of bugs
AFL	105
Driller	121
T-Fuzz	166
Driller - AFL	16
T-Fuzz - AFL	61
T-Fuzz - Driller	55
Driller - T-Fuzz	10

LAVA-M dataset

- T-Fuzz performs well given conditions favorable for VUzzer and Steelix
- T-Fuzz outperforms VUzzer and Steelix for “hard” checks
- T-Fuzz was defeated by Steelix due to transformation explosion in who
- T-Fuzz found 1 unintended bug in who

Program	Total # of bugs	VUzzer	Steelix	T-Fuzz
base64	44	17	43	43
unique	28	27	24	26
md5sum	57	1	28	49
who	2136	50	194	95*

Real-world programs

- Widely evaluated in related work
- T-Fuzz detected far more (verified) crashes than AFL
- T-Fuzz found 3 new bugs

Program + library	AFL	T-Fuzz
pngfix + libpng (1.7.0)	0	11
tiffinfo + libtiff (3.8.2)	53	124
magick + ImageMagick (7.0.7)	0	2
pdftohtml + libpoppler (0.62.0)	0	1

Conclusion & future work

- Fuzzers are limited by coverage and unable to find “deep” bugs
- T-Fuzz extend fuzzing by “mutating” the target program as well
- **Experimental results show that T-Fuzz is more effective than state-of-art fuzzers**
 - ❑ T-Fuzz has improvement over Driller/AFL by 45%/58%
 - ❑ T-Fuzz was able to trigger bugs guarded by “hard” checks
 - ❑ T-Fuzz found new bugs: 1 in LAVA-M dataset and 3 in real world programs
- Future work
 - ❑ Improve transformation strategies
 - ❑ Improve filtering of false positives



<https://github.com/HexHive/T-Fuzz>