




# DUMPLING: FINE-GRAINED DIFFERENTIAL JAVASCRIPT ENGINE FUZZING



# V8 EXECUTION TIERS



# V8 EXECUTION TIERS




Even confusing  $-0.0$  with  $+0.0$  is enough for RCE [Röt18]

# VM STATE

```

0 : 01 0d e8 ff ff 1f LdaSmi.ExtraWide [536870888]
6 : c5 Star0
7 : 13 00 LdaConstant [0]
9 : c2 Star3
10 : 2d f6 01 00 GetNamedProperty r3, [1], [0]
14 : c3 Star2
15 : 5e f7 f6 f9 02 CallProperty1 r2, r3, r0, [2]
20 : c4 Star1
21 : 2d f8 02 04 GetNamedProperty r1, [2], [4]
25 : c3 Star2
26 : 13 03 LdaConstant [3]
28 : c1 Star4
29 : 5f f7 f8 f5 f9 06 CallProperty2 r2, r1, r4, r0, [6]
35 : aa Return

```




# JIT COMPIRATION

```

LdaConstant [0]
Star2
Mov <closure>, r3
CallRuntime [DeclareGlobals], r2 - r3
LdaZero
Star1
LdaUndefined
Star0
LdaSmi.Wide [10000]
TestLessThan r1, [0]
JumpIfFalse [31] (0x4ea00040085 @ 53)
LdaGlobal [1], [1]
Star2
CreateObjectLiteral [2], [3], #41
Star3
Ldar r1
DefineNamedOwnProperty r3, [3], [4]
CallUndefinedReceiver1 r2, r3, [6]
Star0
Ldar r1
Inc [8]
Star1
JumpLoop [34], [0], [9] (0x4ea0004005f @ 15)
Ldar r0
Return

```



```

18 int3l
19 movl rbx,[rcx-0xc]
1c REX.W orq rbx,[r13+0x1e0]
23 testb [rbx+0x1a],0x20
27 jz 0x7fe5e00003b6 B0 <+0x36>
29 REX.W movq r10,0x7fe5b9df2a00 (CompileLazyDeoptimizedCode) ; off heap
33 jmp r10
B0:
36 push rbp
37 REX.W movq rbp,rsp
3a push rsi
3b push rdi
3c push rax
3d REX.W subq rsp,0x8
41 REX.W movq [rbp-0x20],rsi
45 REX.W cmpq rsp,[r13-0x60] (external value (StackGuard::address_of_jslimit)
49 jna 0x7fe5e0000456 B1 <+0xd6>
B2,3:
4f REX.W movq rdx,[rbp+0x18]
53 testb rdx,0x1
56 jz 0x7fe5e0000488 <+0x108>
5c movl rcx,0x298bdd ; (compressed) object: 0xbdd00298bdd <Map[16](HOLE
61 cmpl [rdx-0x1],rcx
64 jnz 0x7fe5e000048c <+0x10c>
6a movl rcx,[rdx+0xb]
6d REX.W movq rdi,0xbdd00284d2d ; object: 0xbdd00284d2d <JSFunction lo
77 movl rsi,[rdi+0x13]
7a REX.W addq rsi,r14
7d push rcx
7e REX.W movq rcx,0xbdd00284c65 ; object: 0xbdd00284c65 <console map =
88 push rcx
89 REX.W leaq rcx,[r14+0x741]
90 push rcx
91 push 0xc
93 push rdi
94 REX.W leaq rax,[r14+0x69]
98 push rax
99 REX.W movq rbx,0x7fe5baad57c0 ; external reference (Builtin_ConsoleLo
a3 movl rax,0x6
a8 REX.W movq rcx,rax
ab REX.W movq r10,0x7fe5ba1885c0 (CEntry_Return1_ArgvOnStack_BuiltinExit)
b5 call r10
b8 REX.W leaq rax,[r14+0x69]
bc REX.W movq rcx,[rbp-0x18]

```

# JIT COMPIILATION




```
LdaConstant [0]
Star2
Mov <closure>, r3
CallRuntime [DeclareGlobals], r2-r3
LdaZero
Star1
LdaUndefined
Star0
LdaSmi.Wide [10000]
TestLessThan r1, [0]
JumpIfFalse [31] (0x4ea00040085 @ 53)
LdaGlobal [1], [1]
Star2
CreateObjectLiteral [2], [3], #41
Star3
Ldar r1
DefineNamedOwnProperty r3, [3], [4]
CallUndefinedReceiver1 r2, r3, [6]
Star0
Ldar r1
Inc [8]
Star1
JumpLoop [34], [0], [9] (0x4ea0004005f @ 15)
Ldar r0
Return
```



Compare VM states from unoptimized execution (left) to optimized execution (right).

```
18 int3l
19 movl rbx,[rcx-0xc]
1c REX.W orq rbx,[r13+0x1e0]
23 testb [rbx+0x1a],0x20
27 jz 0x7fe5e00003b6 B0 <+0x36>
29 REX.W movq r10,0x7fe5b9df2a00 (CompileLazyDeoptimizedCode) ; off heap
33 jmp r10
B0:
36 push rbp
37 REX.W movq rbp,rsp
3a push rsi
3b push rdi
3c push rax
3d REX.W subq rsp,0x8
41 REX.W movq [rbp-0x20],rsi
45 REX.W cmpq rsp,[r13-0x60] (external value (StackGuard::address_of_jslimit
49 jna 0x7fe5e0000456 B1 <+0xd6>
B2,3:
4f REX.W movq rdx,[rbp+0x18]
53 testb rdx,0x1
56 jz 0x7fe5e0000488 <+0x108>
5c movl rcx,0x298bdd ; (compressed) object: 0xbdd00298bdd <Map[16](HOLE
61 cmpl [rdx-0x1],rcx
64 jnz 0x7fe5e000048c <+0x10c>
6a movl rcx,[rdx+0xb]
6d REX.W movq rdi,0xbdd00284d2d ; object: 0xbdd00284d2d <JSFunction local
77 movl rsi,[rdi+0x13]
7a REX.W addq rsi,r14
7d push rcx
7e REX.W movq rcx,0xbdd00284c65 ; object: 0xbdd00284c65 <console map =
88 push rcx
89 REX.W leaq rcx,[r14+0x741]
90 push rcx
91 push 0xc
93 push rdi
94 REX.W leaq rax,[r14+0x69]
98 push rax
99 REX.W movq rbx,0x7fe5baad57c0 ; external reference (Builtin_ConsoleLog
a3 movl rax,0x6
a8 REX.W movq rcx,rax
ab REX.W movq r10,0x7fe5ba1885c0 (CEEntry_Return1_ArgvOnStack_BuiltinExit)
b5 call r10
b8 REX.W leaq rax,[r14+0x69]
bc REX.W movq rcx,[rbp-0x18]
```

# OVERVIEW



## TRACES

Execution traces  
even during JIT

## MATCHING

Matching algorithm  
to compare traces


## DUMPLING

Differential Fuzzer  
using our bug oracle

## V8 BUGS

Evaluation and 8  
new V8 bugs

# STATE EXTRACTION



# STATE EXTRACTION: JIT

- State is spread accross machine registers and stack

```
B0:  
29  push rbp  
2a  REX.W movq rbp,rsp  
2d  push rsi  
2e  push rdi  
2f  push rax  
30  REX.W subq rsp,0x8  
34  REX.W movq [rbp-0x20],rsi  
38  REX.W cmpq rsp,[r13-0x60] (external value (StackGuard::address_of_jslimit()))  
3c  jna 0x7f8d89f84134 B1,14 <+0xf4>  
  
B2,3:  
42  REX.W movq rcx,[rbp+0x18]  
46  testb rcx,0x1  
49  jz 0x7f8d89f841aa <+0x16a>  
  
54  movl rdi,0x99e75 ; (compressed) object: 0x28ba00099e75 <Map[16](HOLEY_ELEMENTS)>  
59  cmpl [rcx-0x1],rdi  
5c  jnz 0x7f8d89f841ae <+0x16e>  
  
67  movl r8,[rcx+0xb]  
6b  REX.W movq r9,[rbp+0x20]  
6f  testb r9,0x1  
73  jz 0x7f8d89f841b2 <+0x172>  
  
7e  cmpl [r9-0x1],rdi  
82  jnz 0x7f8d89f841b6 <+0x176>
```

# STATE EXTRACTION: JIT

- State is spread accross machine registers and stack
- **How** do we get back to state comparable to interpreter execution?
- **Where** is state extraction possible?

```
B0:  
29  push rbp  
2a  REX.W movq rbp,rsp  
2d  push rsi  
2e  push rdi  
2f  push rax  
30  REX.W subq rsp,0x8  
34  REX.W movq [rbp-0x20],rsi  
38  REX.W cmpq rsp,[r13-0x60] (external value (StackGuard::address_of_jslimit()))  
3c  jna 0x7f8d89f84134 B1,14 <+0xf4>  
  
B2,3:  
42  REX.W movq rcx,[rbp+0x18]  
46  testb rcx,0x1  
49  jz 0x7f8d89f841aa <+0x16a>  
  
54  movl rdi,0x99e75 ; (compressed) object: 0x28ba00099e75 <Map[16](HOLEY_ELEMENTS)>  
59  cmpl [rcx-0x1],rdi  
5c  jnz 0x7f8d89f841ae <+0x16e>  
  
67  movl r8,[rcx+0xb]  
6b  REX.W movq r9,[rbp+0x20]  
6f  testb r9,0x1  
73  jz 0x7f8d89f841b2 <+0x172>  
  
7e  cmpl [r9-0x1],rdi  
82  jnz 0x7f8d89f841b6 <+0x176>
```

# STATE EXTRACTION: JIT

- State is spread accross machine registers and stack
- **How** do we get back to state comparable to interpreter execution?
- **Where** is state extraction possible?
- No influence on JS execution semantics and JIT compiler optimizations

```
B0:  
29  push rbp  
2a  REX.W movq rbp,rsp  
2d  push rsi  
2e  push rdi  
2f  push rax  
30  REX.W subq rsp,0x8  
34  REX.W movq [rbp-0x20],rsi  
38  REX.W cmpq rsp,[r13-0x60] (external value (StackGuard::address_of_jslimit()))  
3c  jna 0x7f8d89f84134 B1,14 <+0xf4>  
  
B2,3:  
42  REX.W movq rcx,[rbp+0x18]  
46  testb rcx,0x1  
49  jz 0x7f8d89f841aa <+0x16a>  
  
54  movl rdi,0x99e75 ; (compressed) object: 0x28ba00099e75 <Map[16](HOLEY_ELEMENTS)>  
59  cmpl [rcx-0x1],rdi  
5c  jnz 0x7f8d89f841ae <+0x16e>  
  
67  movl r8,[rcx+0xb]  
6b  REX.W movq r9,[rbp+0x20]  
6f  testb r9,0x1  
73  jz 0x7f8d89f841b2 <+0x172>  
  
7e  cmpl [r9-0x1],rdi  
82  jnz 0x7f8d89f841b6 <+0x176>
```

# DEOPTIMIZATION POINTS

```
function f(o1, o2) {  
    return o1.a * o2.a;  
}
```

- Deopt points guard usage of speculative assumption
- JIT tracks context to restore VM state at deopt points

```
B0:  
29 push rbp  
2a REX.W movq rbp,rsp  
2d push rsi  
2e push rdi  
2f push rax  
30 REX.W subq rsp,0x8  
34 REX.W movq [rbp-0x20],rsi  
38 REX.W cmpq rsp,[r13-0x60] (external value (StackGuard::address_of_jslimit()))  
3c jna 0x7f8d89f84134 B1,14 <+0xf4>  
  
B2,3:  
42 REX.W movq rcx,[rbp+0x18]  
46 testb rcx,0x1  
49 jz 0x7f8d89f841aa <+0x16a>  
  
54 movl rdi,0x99e75 ; (compressed) object: 0x28ba00099e75 <Map[16](HOLEY_ELEMENTS)>  
59 cmpl [rcx-0x1],rdi  
5c jnz 0x7f8d89f841ae <+0x16e>  
  
67 movl r8,[rcx+0xb]  
6b REX.W movq r9,[rbp+0x20]  
6f testb r9,0x1  
73 jz 0x7f8d89f841b2 <+0x172>  
  
7e cmpl [r9-0x1],rdi  
82 jnz 0x7f8d89f841b6 <+0x176>
```

46 testb rcx,0x1  
49 jz 0x7f8d89f841aa <+0x16a>

IsObject(o1)

# DEOPTIMIZATION POINTS

```
function f(o1, o2) {  
    return o1.a * o2.a;  
}
```

- Deopt points guard usage of speculative assumption
- JIT tracks context to restore VM state at deopt points
- → Deopt points as natural probing positions for interesting state

```
B0:  
29 push rbp  
2a REX.W movq rbp, rsp  
2d push rsi  
2e push rdi  
2f push rax  
30 REX.W subq rsp, 0x8  
34 REX.W movq [rbp-0x20], rsi  
38 REX.W cmpq rsp, [r13-0x60] (external value (StackGuard::address_of_jslimit()))  
3c jna 0x7f8d89f84134 B1,14 <+0xf4>  
  
B2,3:  
42 REX.W movq rcx, [rbp+0x18]  
46 testb rcx, 0x1  
49 jz 0x7f8d89f841aa <+0x16a>  
  
54 movl rdi, 0x99e75 ; (compressed) object: 0x28ba00099e75 <Map[16](HOLEY_ELEMENTS)>  
59 cmpl [rcx-0x1], rdi  
5c jnz 0x7f8d89f841ae <+0x16e>  
  
67 movl r8, [rcx+0xb]  
6b REX.W movq r9, [rbp+0x20]  
6f testb r9, 0x1  
73 jz 0x7f8d89f841b2 <+0x172>  
  
7e cmpl [r9-0x1], rdi  
82 jnz 0x7f8d89f841b6 <+0x176>
```

46 testb rcx, 0x1  
49 jz 0x7f8d89f841aa <+0x16a>

IsObject(o1)

# DUMPING DURING SPECULATIVE JIT EXECUTION

1. Save state
2. Build VM state
3. Rematerialize escaped values
4. "Dump" VM state
5. Restore state and continue JIT execution

→ partial use of existing deopt mechanism

```
B0:  
29 push rbp  
2a REX.W movq rbp, rsp  
2d push rsi  
2e push rdi  
2f push rax  
30 REX.W subq rsp, 0x8  
34 REX.W movq [rbp-0x20], rsi  
38 REX.W cmpq rsp, [r13-0x60] (external value (StackGuard::address_of_jslimit()))  
3c jna 0x7f8d89f84134 B1,14 <+0xf4>  
  
B2,3:  
42 REX.W movq rcx, [rbp+0x18]  
46 testb rcx, 0x1  
49 jz 0x7f8d89f841aa <+0x16a>  
4f call 0x7f8d29fa1c0 (DumpTurboFrame) ;; near builtin entry  
54 movl rdi, 0x99e75 ;; (compressed) object: 0x28ba00099e75 <Map[16](HOLEY_ELEMENTS)>  
59 cmpl [rcx-0x1], rdi  
5c jnz 0x7f8d89f841ae <+0x16e>  
62 call 0x7f8d29fa1c0 (DumpTurboFrame) ;; near builtin entry  
67 movl r8, [rcx+0xb]  
6b REX.W movq r9, [rbp+0x20]  
6f testb r9, 0x1  
73 jz 0x7f8d89f841b2 <+0x172>  
79 call 0x7f8d29fa1c0 (DumpTurboFrame) ;; near builtin entry  
7e cmpl [r9-0x1], rdi  
82 jnz 0x7f8d89f841b6 <+0x176>  
88 call 0x7f8d29fa1c0 (DumpTurboFrame) ;; near builtin entry
```

**Call Dumpling hook if deopt point not hit**

# **STATE EXTRACTION: DUMPLING MODE - INTERPRETER**


- Optimized run reports dump locations to the fuzzer
- Hook bytecode execution and extract state at those dump locations

# STATE SERIALIZATION

```
-----TurboFan frame dump-----
pc: 7
acc: 13.37
a0: <Object>{
  __proto__: <Class C7>{<String[1]: f>[enumerable]<JSArray>[],
  <String[1]: a>[configurable][enumerable]42(enum cache: 2),
  <String[1]: f>[configurable][enumerable]13.37(enum cache: 0)
}
r0: -INFINITY
context: <ScriptContext[4]>
receiver: <JSGlobalProxy>
closure: <JSFunction f0>
Function ID: 27
```


- Invariant across execution tiers
- Fine-grained and in-depth
- Concise to minimize transmission overhead

# DIFFERENTIAL ORACLE



- No 1:1 mapping of dumps
- Any JIT dump must have an interpreter equivalent in the **same** function invocation

# EVALUATION: OVERHEAD



Fuzzer	Fuzzilli	JIT-Picker	FuzzJIT	Dumpling
Executions	63,775,062	99,240,042	61,434,736	51,535,553

# BUGS

Found 8 new V8 bugs 

Bug Id	Kind	Status	Changes	By	Description
CR41488094	Diff	fixed	+28/-23	D, J	Enumerating properties eagerly, has incorrect side effect
CR335310000	Diff	fixed	+15/0	D	Property not marked as enumerable by Maglev and TurboFan
CR332745405	Diff	fixed	+5/0	D	DefineOwnProperty called the setter of an existing accessor property
CR329330868	assert	dup	N/A	D, J	array.shift does not update pointers in spill slots
CR41484971	Diff	fixed	+52/-40	D	Store inline cache handlers are incorrectly used for defining properties
V8:14605	Diff	fixed	+1/-1	D	The JumpLoop bytecode does not clobber the accumulator in all cases
CR345960102	Diff	fixed	+6/-4	D	TurboFan incorrectly optimizes 64 bit BigInt shifts
CR346086168	Diff	fixed	+109/-107	D	Overflow in BigInt64 shift optimization
V8:14556	Diff	available	N/A	D	The arguments array is handled differently in optimizing compilers
CR40945996	assert	dup	N/A	D	The profiler in Maglev interferes with deoptimization

# CASE STUDY

```
function A() {
  Object.defineProperty(this, "x", { writable: true, configurable: true, value: undefined });
}

class B extends A {
  x = {};
}

for (let i = 0; i < 100; i++) {
  new B();
}
```

Here not "visible", but already detectable by Dumpling

# CASE STUDY

```
function A() {
    Object.defineProperty(this, "x", { writable: true, configurable: true, value: undefined });
}

class B extends A {
    x = {};
}

for (let i = 0; i < 100; i++) {
    new B();
}
```

Here not "visible", but already detectable by Dumpling

Other fuzzers need generate something like

```
let b = new B();
console.log(b.propertyIsEnumerable("x"));
```

optimizations enabled: "true", optimizations disabled: "false"

# **CONCLUSION**

## **KEY PROBLEM**

Find differentials between JS engine execution tiers automatically

## **DUMPLING**

Extract VM states during runtime and compare between JIT and interpreter

Leveraging deoptimization points, a mechanism already implemented in JS engines

## **RESULT**

Find bugs before they become "visible"

# QUESTIONS?

Find our artifact here: [github.com/two-heart/dumpling-artifact-evaluation](https://github.com/two-heart/dumpling-artifact-evaluation)



('@95p@mastodon.cloud)

(@NearBeteigeuze)

(liam@seine.email)

# BIBLIOGRAPHY

- [GSV22] Jakob Gruber, Leszek Swirski, and Toon Verwaest. Maglev. 2022. url: [https://docs.google.com/document/d/13CwgSL4yawxuYg3iNIM-4ZPCB8RgJya6b8H\\_E2F-Aek/](https://docs.google.com/document/d/13CwgSL4yawxuYg3iNIM-4ZPCB8RgJya6b8H_E2F-Aek/) (visited on 11/28/2023).
- [Röt18] Stephen Röttger. Chrome: V8: incorrect type information on Math.expm1. 2018. url: <https://crbug.com/project-zero/1710> (visited on 03/18/2024).
- [Flü16] Olivier Flückiger. Ignition: V8 Interpreter. 2016. url: <https://docs.google.com/document/d/11T2CRex9hXxoJwbYqVQ32yIPMh0ouUZLdyrtmMoL44> (visited on 11/20/2023).

